

## COSC 6397 Big Data Analytics

### Master Worker Programming Pattern

Edgar Gabriel  
Spring 2015

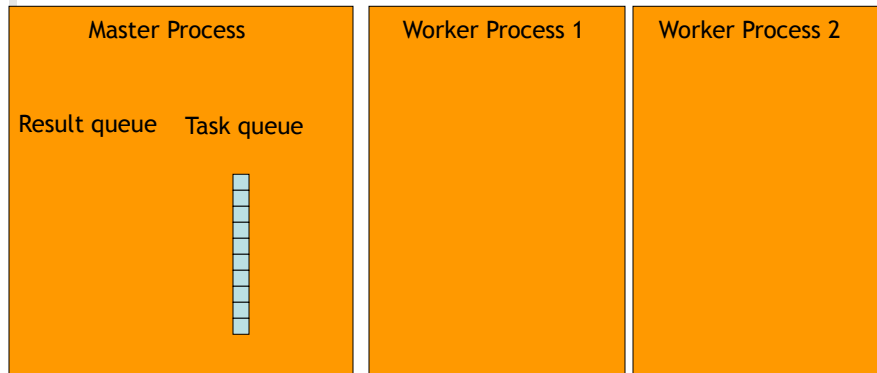
UNIVERSITY of **HOUSTON**

### Master-Worker pattern

- General idea: distribute the work among a number of processes
- Two logically different entities:
  - master process: manages assignments to worker processes
  - worker processes executing a task assigned to them by the master
- Communication only occurs between master process and worker processes
  - No direct worker to worker communication possible

UNIVERSITY of **HOUSTON**

## Master-Worker framework



UNIVERSITY of **HOUSTON**

## Master process pseudo code

```

for each worker w=0...no. of workers {
    f = FETCH next piece of work
    SEND f to w
}
while ( not done) {
    RECEIVE result from any process
    w = process who sent result
    STORE result
    f = FETCH next piece of work
    if f != no more work left
        SEND f to w
    else
        SEND termination message to w
}
    
```

UNIVERSITY of **HOUSTON**

## Master process

- List of things to worry about:
  - Number of workers:
    - Want to maximize performance -> many workers
    - Too many workers: master process can become the bottleneck
  - Work distribution: too little work per worker can lead to large management overhead for the master
  - SEND/RECEIVE
    - can be implemented using any communication paradigm (sockets, MPI, http, files on a shared filesystem,...)
    - More than one message might be required per communication step

UNIVERSITY of **HOUSTON**

## Master process

- Support for process failure
  - Master needs to maintain a list of which work item has been assigned to which worker
  - If a worker fails, work is reassigned to new worker
  - Resilience of master can be achieved through reliable back-end storage
- Correctness verification :
  - assign the same work to more than one worker
  - compare results obtained

UNIVERSITY of **HOUSTON**

## Worker process pseudo code

```
while ( not done ) {  
    RECEIVE work f from master  
    If f equals termination signal  
        exit  
    Else  
        result r = execute work on f  
        SEND r to master  
}
```

UNIVERSITY of **HOUSTON**

## Worker process

- Dynamically generating worker instances vs. using a fixed number of workers
- Worker can support multiple functions
  - Additional message might be required to identify the task that the worker needs to execute

UNIVERSITY of **HOUSTON**

## 1<sup>st</sup> Example: Word Count

- Application counting the number of occurrences of each word in a given input file
- **Input:** a file with n number of lines
- **Output:** list of words and the number of occurrences
- Two step approach:
  - Step 1:
    - Each worker gets one/a fixed no. of line(s) of the file
    - Worker returns a list of <word, #occurrences> to master.
    - Master writes list of words, #occurrences into a temporary file
  - Step 2: the same word will appear multiple times in the temporary file
    - Sort temporary file
    - Add the number of occurrences of the same word
    - Write word and final number of occurrences to end file

## Word count step 1

```

for each worker w=0...no. of workers {
    f = READLINE (file)
    SEND f to w
}
while ( not done) {
    RECEIVE number of elements in the list returned
    allocate temporary buffer to receive list
    RECEIVE list <words, occurrences>
    WRITE list of <words, occurrences> to temporary file
    release temporary buffer
    f = READLINE (file)
    if f != EOF      SEND f to w
    else            SEND termination message to w
}
  
```

## Word count step 2

```

SORT temporary file based on word
<word, occurrence> = READ (temporary file);
currentword = word;
currentcount = occurrence;
while ( not done) {
    <word, occurrence> = READ (temporary file);
    if ( word != currentword ) {
        WRITE (outputfile, currentword, currentcount)
        currentword = word;
        currentcount = 0;
    }
    currentcount += occurrence;
}

```

UNIVERSITY of HOUSTON

## Alternative design possibilities

- Master searches for word occurrence and adds value received from worker process immediately
  - Avoids temporary file
  - Requires large number of search operations in the file
  - Only final output file is sorted
- Worker process does not return list of <word, occurrence> for each line received
  - Just signals that its ready for a new assignment
  - returns only final list after termination signal

UNIVERSITY of HOUSTON

## Alternative design possibilities (II)

- Using workers for step 2 often also necessary
- Using workers for sorting
  - Requires sending large data volumes
  - Requires direct worker to worker communication
  - > virtually impossible in a classic master-worker setting
  - > Amdahl's law strikes!
- Aggregation of multiple entries for the same word can be distributed to workers

UNIVERSITY of **HOUSTON**

## 2<sup>nd</sup> Example: k-means clustering

- An unsupervised clustering algorithm
- “K” stands for number of clusters, typically a user input to the algorithm
- Iterative algorithm
  - Might obtain only a local minimum
  - Works only for numerical data
- $x_1, \dots, x_N$  are data points
- Each data point (vector  $x_i$ ) will be assigned to exactly one cluster
- Dissimilarity measure: Euclidean distance metric

UNIVERSITY of **HOUSTON**

## K-means Algorithm

- $C(i)$ : cluster that data point  $x_i$  is currently assigned to
- For a current set of cluster means  $m_k$ , assign each data point to exactly one cluster:

$$C(i) = \arg \min_{1 \leq k \leq K} \|x_i - m_k\|^2, i = 1, \dots, N$$

- $N_k$ : number of data points in cluster  $k$
- For a given cluster assignment  $C$  of the data points, compute the cluster means  $m_k$ :

$$m_k = \frac{\sum_{i: C(i)=k} x_i}{N_k}, k = 1, \dots, K.$$

- Iterate above two steps until convergence

UNIVERSITY of **HOUSTON**

## K-means clustering: sequential version

Input :

$X = \{x_1, \dots, x_k\}$  Instances to be clustered

$K$  : Number of clusters

Output

$M = \{m_1, \dots, m_k\}$  : cluster centroids

$m: X \rightarrow C$  : cluster membership

Algorithm

Set initial value for  $M$

while  $m$  has changed

$c_n^{temp} = 0, n = 1, \dots, k$

$N_n = 0, n = 1, \dots, k$

for each  $x \in X$

$C(j) = \min \text{distance}(x_j, c_n), n = 1, \dots, k$

$c_n^{temp} += x_j$

$N_n ++$

end

recompute  $C$  based on  $c_n^{temp}$  and  $N$

**STON**



## K-means clustering

- Initial cluster means
  - Using some other/simpler pre-clustering algorithms
  - Random values
- Random values difficult when using multiple processes
  - Pseudo random number generators
  - Master process creates random numbers and distributes it to workers

UNIVERSITY of HOUSTON

## Master-worker k-means

```

m= randomly generated initial cluster means
x = data points
for each worker w=0...no. of workers {
    SEND portion of x for worker w
}

Do {
    for each worker w=0...no. of workers {
        SEND m to w
    }
    for each worker w=0...no. of workers {
        RECEIVE C from w
    }
    Calculate N and ctemp for all cluster
    Recalculate m
} while (m has changed )
  
```

UNIVERSITY of HOUSTON

## Master worker k-means

- Portion of data points  $x$  assigned to a worker remains constant -> has to be sent only once
- Need to separate loops to send  $m$  and receive  $C$  to avoid serialization of the worker processes
- Problem is strongly synchronized
  - All worker processes have to finish before master process can recalculate  $N_k$ ,  $ctemp$  and  $m$
  - Calculating  $N_k$ ,  $ctemp$  and  $m$  sequential in this version
  - > Amdahl's law strikes!

UNIVERSITY of **HOUSTON**

## Summary

- Master worker pattern useful for a number of application scenarios
  - Mostly 'trivially' parallel problems
- Supports dynamic load balancing
  - A more powerful worker will return work faster and get more work assigned than a slow worker
  - Well suited for heterogeneous environments
  - Easy to integrate fault tolerance
- Not well suited for synchronized problems
- Amdahl's law will prevent the model from scaling up to extreme problems

UNIVERSITY of **HOUSTON**